

# R Homework Three

Katherine Wolf

Introduction to Causal Inference (PH252D)

April 7, 2020

## 1 Background story.

## 2 Import and explore the data set RAssign3.csv.

### 2.1 Use the read.csv function to import the dataset and assign it to dataframe obs\_data.

```
library(tidyverse)

obs_data <- read_csv("RAssign3.csv")
```

### 2.2 Use the names, tail, and summary functions to explore the data.

```
names(obs_data)

## [1] "W1" "W2" "W3" "W4" "W5" "Y"

tail(obs_data)

## # A tibble: 6 x 6
##   W1     W2     W3     W4     W5     Y
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     0     0 0.609 0.393     2  92.5
## 2     1     0 1.47  0.713     3  95.7
## 3     0     0 0.0843 0.448     2  90.4
## 4     0     0 1.13  0.160     4  97.5
## 5     0     1 0.207 0.444     1 112.
## 6     0     0 0.435 0.121     1  99.0

summary(obs_data)

##           W1           W2           W3           W4
##  Min.    :0.000   Min.    :0.0000   Min.    :0.000327   Min.    :0.1192
## 1st Qu.:0.000   1st Qu.:0.0000   1st Qu.:1.194527   1st Qu.:0.2839
##  Median :0.000   Median :1.0000   Median :2.486190   Median :0.3980
##   Mean   :0.108   Mean   :0.5082   Mean   :2.474991   Mean   :0.4174
## 3rd Qu.:0.000   3rd Qu.:1.0000   3rd Qu.:3.731702   3rd Qu.:0.5522
##   Max.   :1.000   Max.   :1.0000   Max.   :4.998937   Max.   :0.8807
##           W5           Y
##  Min.    :1.000   Min.    : 88.00
## 1st Qu.:1.000   1st Qu.: 99.28
##  Median :2.000   Median :109.39
```

```
## Mean      :1.897      Mean      :109.39
## 3rd Qu.   :2.000      3rd Qu.   :119.30
## Max.      :4.000      Max.      :137.65
```

**2.3 Use the `nrow` function to count the number of communities in the data set. Assign this number as `n`.**

```
n <- nrow(obs_data)

n

## [1] 5000
```

### 3 Code discrete Super Learner to select the estimator with the lowest cross-validated risk estimate.

**3.1 Briefly discuss the motivation for using discrete Super Learner (a.k.a. the cross-validation selector).**

*add explanation here*

**3.2 Create the following transformed variables and add them to the data frame `obs_data`:**

- `sin_W3 <- sin(obs_data$W3)`
- `W4_sq <- obs_data$W4 * obs_data$W4`
- `cos_W5 <- cos(obs_data$W5)`

```
obs_data$sin_W3 <- sin(obs_data$W3)
obs_data$W4_sq <- obs_data$W4 * obs_data$W4
obs_data$cos_W5 <- cos(obs_data$W5)
```

**3.3 Split the data into  $V = 20$  folds. Create the vector `fold` and add it to the data frame `obs_data`.**

```
# With n = 5000 observations total, we want n/20 = 250 observations in each fold.

obs_data$fold <- c(rep(1, 250),
                  rep(2, 250),
                  rep(3, 250),
                  rep(4, 250),
                  rep(5, 250),
                  rep(6, 250),
                  rep(7, 250),
                  rep(8, 250),
                  rep(9, 250),
                  rep(10, 250),
                  rep(11, 250),
                  rep(12, 250),
                  rep(13, 250),
                  rep(14, 250),
```

```
rep(15, 250),
rep(16, 250),
rep(17, 250),
rep(18, 250),
rep(19, 250),
rep(20, 250))
```

**3.4 Create an empty matrix `CV_risk` with 20 rows and 4 columns for each algorithm, evaluated at each fold.**

```
cv_risk <- matrix(NA, nrow=20, ncol=4)
```

**3.5 Use a for loop to fit each estimator on the training set (19/20 of the data); predict the expected MUAC for the communities in the validation set (1/20 of the data), and evaluate the cross-validated risk.**

1. Since each fold needs to serve as the training set, have the for loop run from V is 1 to 20.
2. Create the validation set as a data frame `valid`, consisting of observations with fold equal to V.
3. Create the training set as a data frame `train`, consisting of observations with fold not equal to V.
4. Use `glm` to fit each algorithm on the training set. Be sure to specify `data = train`.
5. For each algorithm, predict the average MUAC for each community in the validation set. Be sure to specify the `type = 'response'` and `newdata = valid`.
6. Estimate the cross-validated risk for each algorithm with the L2 loss function. Take the mean of the squared differences between the observed outcomes Y in the validation set and the predicted outcomes. Assign the cross-validated risks as a row in the matrix `cv_risk`.

```
for (V in 1:20){
  valid <- obs_data[obs_data$fold == V,]
  train <- obs_data[obs_data$fold != V,]
  model_a <- glm(Y ~ W1 + W2 + sin_W3 + W4_sq, data = train)
  model_b <- glm(Y ~ W1 + W2 + W4 + cos_W5, data = train)
  model_c <- glm(Y ~ W2 + W3 + W5 + W2:W5 + W4_sq + cos_W5, data = train)
  model_d <- glm(Y ~ W1*W2*W5, data = train)

  predict_a <- predict(object = model_a, type = 'response', newdata = valid)
  predict_b <- predict(object = model_b, type = 'response', newdata = valid)
  predict_c <- predict(object = model_c, type = 'response', newdata = valid)
  predict_d <- predict(object = model_d, type = 'response', newdata = valid)

  cv_risk[V,1] <- mean((valid$Y - predict_a)^2)
  cv_risk[V,2] <- mean((valid$Y - predict_b)^2)
  cv_risk[V,3] <- mean((valid$Y - predict_c)^2)
  cv_risk[V,4] <- mean((valid$Y - predict_d)^2)
}
```

**3.6 Select the algorithm with the lowest average cross-validated risk. Hint: Use the `colMeans` function.**

```

# get the average risks
average_risks <- colMeans(cv_risk)

# restore their model names
names(average_risks) <- c("model_a", "model_b", "model_c", "model_d")

# print them
average_risks

##   model_a   model_b   model_c   model_d
## 8.312289 13.025325  7.762348 16.031240

# find the average risk that minimizes the L2 loss function
minimum_average_risk <- average_risks[average_risks == min(average_risks)]

# print it
minimum_average_risk

##   model_c
## 7.762348

# get the model name
best_discrete_model_name <- names(minimum_average_risk)

```

### 3.7 Fit the chosen algorithm on all the data.

```

# fit the best model
# (Note: This code takes the name of whatever model came out on top
# as a character string stored in the variable best_discrete_model_name,
# deparse-substitutes it back to a variable to get the
# glm() model object, pulls its formula using formula(), and then
# uses that to run the new glm(), replacing
# the data with the full dataset obs_data. I did that so that
# I wouldn't have to check manually which model won in the prior
# step.)
discrete_best_estimate <-
  glm(formula = formula(deparse(substitute(best_discrete_model_name))),
      data = obs_data)

discrete_best_estimate

##
## Call:  glm(formula = formula(deparse(substitute(best_discrete_model_name))),
##      data = obs_data)
##
## Coefficients:
## (Intercept)          W2          W3          W5          W4_sq        cos_W5
##   95.64175    12.84349    2.02159    0.07781   -10.80041    2.11408
##      W2:W5
##   4.87470
##
## Degrees of Freedom: 4999 Total (i.e. Null);  4993 Residual
## Null Deviance:      618800

```

```
## Residual Deviance: 38700 AIC: 24440

summary(discrete_best_estimate)

##
## Call:
## glm(formula = formula(deparse(substitute(best_discrete_model_name))),
##      data = obs_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.6682  -1.9842  -0.4801   0.9096   9.5123
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  95.64175    0.26641  359.001  <2e-16 ***
## W2           12.84349    0.21116   60.824  <2e-16 ***
## W3            2.02159    0.02716   74.424  <2e-16 ***
## W5            0.07781    0.14745    0.528    0.598
## W4_sq       -10.80041    0.29801  -36.242  <2e-16 ***
## cos_W5        2.11408    0.18974   11.142  <2e-16 ***
## W2:W5         4.87470    0.09881   49.335  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 7.75025)
##
##      Null deviance: 618826  on 4999  degrees of freedom
## Residual deviance:  38697  on 4993  degrees of freedom
## AIC: 24437
##
## Number of Fisher Scoring iterations: 2
```

### 3.8 Can we do better?

I bet that we can, since we just arbitrarily picked four models and then stipulated that we had to choose one and only one of them.

## 4 Use the SuperLearner package to build the best combination of algorithms.

### 4.1 Load the Super Learner package with the library function and set the seed to 252.

```
library(SuperLearner)

## Warning: package 'SuperLearner' was built under R version 3.6.3
## Loading required package: nnls
## Super Learner
## Version: 2.0-26
## Package created on 2019-10-27

set.seed(252)
```

- 4.2 Use the `source` function to load script file `Rassign3.Wrappers.R`, which includes the wrapper functions for the *a priori* specified parametric regressions.

```
source("Rassign3.Wrappers.R")
```

- 4.3 Specify the algorithms to be included in Super Learner's library.

```
sl_library <- c('SL.glm.EstA',  
               'SL.glm.EstB',  
               'SL.glm.EstC',  
               'SL.glm.EstD',  
               'SL.ridge',  
               'SL.rpartPrune',  
               'SL.polymars',  
               'SL.mean')
```

*Bonus: Very briefly describe the algorithms corresponding to `SL.ridge`, `SL.rpartPrune`, `SL.polymars` and `SL.mean`.*

do this at the end

- 4.4 Create data frame `X` with the predictor variables.

Include the original predictor variables and the transformed variables.

- 4.5 Run the `SuperLearner` function. Be sure to specify the outcome `Y`, the predictors `X`, and the library `SL.library`. Also include `cvControl=list(V=20)` in order to get 20-fold cross-validation.

- 4.6 Explain the output to relevant policy makers and stake-holders. What do the columns `Risk` and `Coef` mean? Are the cross-validated risks from `SuperLearner` close to those obtained by your code?

do this at the end

## 5 Implement `CV.SuperLearner`.

- 5.1 Explain why we need `CV.SuperLearner`.

do this at the end

- 5.2 Run `CV.SuperLearner`.

- 5.3 Explore the output. Only include the output from the `summary` function in your write-up, but comment on the other output.

## 6 Bonus!

- 6.1 Try adding more algorithms to the `SuperLearner` library.

do this at the end

- 6.2 Try writing your own wrapper function.

do this at the end